

SATHAPPAN MUTHIAH

3919, 4th St N, Apt 4,
Arlington, VA-22203

Github: <https://github.com/sathappanspm>

Work Permit: F1 until 2020

E-mail: sathap1@vt.edu

Phone: 801-201-0619

Website: <http://people.cs.vt.edu/sathap1/>

SUMMARY

SUMMARY

I am a Phd student working in spatio-temporal event forecasting using open-source indicators (OSI). Specifically my work focuses on methods for 1) identifying and detecting planned events from text, 2) automated geo-coding approaches for text data (tweet, news etc), 3) fusing and suppressing alerts from multiple models with dynamic thresholding to control precision and recall and 4) building cost-effective methods to encode events in text into tabular form.

Experienced with working on big data, time series and text mining. Broadly interested in forecasting, text mining / Applied NLP, machine learning, and data science.

Projected to graduate in December 2020

EDUCATION

- **PhD** in Computer Science June 2014 — present
Virginia Tech
Arlington, VA
GPA: 3.80
Advisor: Dr. Naren Ramakrishnan
- **MS.** in Computer Science Jan 2012 — June 2014
Virginia Tech
Arlington, VA
GPA: 3.82
Advisor: Dr. Naren Ramakrishnan
- **B.Tech. in Information Tech.** 2007 — 2011
National Institute of Tech.
Bhopal, India
GPA: 8.65

PROFESSIONAL EXPERIENCE

- **Virginia Tech** Spring 2012 — Present
GRA, Discovery Analytics Center
Arlington, VA
 - **Time Series Forecasting for Cyber-Physical Systems:** Studied utility of Sequence-to-Sequence (Seq2Seq) models for the task of multi-variate time series state forecasting. Characterized need for changes to traditional seq2seq architecture (developed predominantly for NLP) for applying in time-series and thereby collaborated in effort to build dynamic attention networks (DyAT) for state forecasting. The work is currently under review in IJCAI 2019
 - **Human-aware Approaches for Event Encoding:** The project is aimed building a cost effective system for encoding political events like Protests, Military Actions etc from news media. The major goal of the system is to create a system with better understanding of uncertainty i.e., it knows when it doesn't know. Characterizing uncertainty helps determine which sections/documents need human supervision. Overall the goal of the project is to maximize coverage via automated approaches for a given accuracy level and at the same time identify documents that can be sent for human supervision (with max available time constraints) to improve overall performance of the hybrid AI + Human system.
 - **EMBERS (IARPA Open Source Indicators program):** Worked as a part of multi-disciplinary multi-university team on building a real-time forecasting system of spatio-temporal events like Civil Unrest. EMBERS was sponsored by a three-year contract for approx. \$13.36M from the IARPA OSI Program. Led development of model to identify planned events from social media, news and other sources and approaches to create automated narrative (sequence of) stories from real-time news. Also researched on approaches to perform automated geo-coding (geo-disambiguation) of news and other textual content. Besides contributed in system architecture and ingestion/enrichment of data sources.
 - **SAFE (IARPA Mercury):** Led development of fusion system. The fusion system is responsible for accepting alerts from all underlying models and deciding which alerts (each alert is a tuple of <forecast date, location, actors, event-type>) to merge and which ones to suppress. Built models to perform real-time clustering of alerts and to predict expected quality of an incoming alert which can then be used to merge/suppress alerts.

- **IGACAT (IARPA Functional Genomic and Computational Assessment of Threats Program):** Researching on methods to characterize threats from specific genes/proteins using heterogeneous interaction graph created from PubMed/Medline abstracts and information of GO functions, sequence groups and cluster information obtained from knowledge bases like UNIPROT, TrEMBL etc. Built deep learning models (CNN, LSTM) for performing few shot learning to identify functions of a nucleotide sequence by embedding both the sequence and label (the functions) in the same space.

- **Indian Institute of Scientific Education and Research**

Intern, advised by Dr. Prashanta Panigrahi

Nov 2008 — Jan 2009

Kolkata, India

- Studied the feasibility of a quantum computing system using common and attainable finite dimensional multipartite quantum states.

- **National Institute of Oceanography**

Intern, mentored by Dr. Biswajit Chakraborty

May-June 2010

Goa, India

- Estimated geophysical parameters using SONAR Backscatter data obtained off an experiment on the Western Continental shelf of India. Developed processing and cleansing techniques to process backscatter data and used simplex method for estimation.

PUBLICATIONS

Accepted

N. Muralidhar, S. Muthiah, K. Nakayama, N. Ramakrishnan, and R. Sharma. Multivariate long-term state forecasting in cyber-physical systems: A sequence to sequence approach. *Big Data*, 2019

M. R. Islam, S. Muthiah, and N. Ramakrishnan. NActSeer: Predicting User Actions in Social Network using Graph Augmented Recurrent Neural Network. In *Accepted at CIKM 2019*, 2019

M. R. Islam, S. Muthiah, and N. Ramakrishnan. RumorSleuth: Joint Detection of Rumor Veracity and User Stance. In *2019 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining (ASONAM)*. IEEE, 2019

N. Muralidhar, S. Muthiah, and N. Ramakrishnan. Dyat nets: Dynamic attention networks for state forecasting in cyber-physical systems. In *Proceedings of the Twenty-Eighth International Joint Conference on Artificial Intelligence, IJCAI-19*, pages 3180–3186. International Joint Conferences on Artificial Intelligence Organization, 7 2019

D. K. Gupta, S. Muthiah, D. Mares, and N. Ramakrishnan. Forecasting Civil Strife: An Emerging Methodology. In *HUSO The Third International Conference on Human and Social Analytics*, 2017

P. Chakraborty, S. Muthiah, R. Tandon, and N. Ramakrishnan. Hierarchical Quickest Change Detection via Surrogates. *arXiv preprint arXiv:1603.09739*, 2016

S. Muthiah, B. Huang, J. Arredondo, D. Mares, L. Getoor, G. Katz, and N. Ramakrishnan. Capturing planned protests from open source indicators. *AI Magazine*, 37(2), 2016

S. Muthiah, P. Butler, R. P. Khandpur, P. Saraf, N. Self, A. Rozovskaya, L. Zhao, J. Cadena, C.-T. Lu, A. Vullikanti, A. Marathe, K. Summers, G. Katz, A. Doyle, J. Arredondo, D. K. Gupta, D. Mares, and N. Ramakrishnan. Embers at 4 years: Experiences operating an open source indicators forecasting system. In *Proceedings of the 22Nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '16, pages 205–214, New York, NY, USA, 2016. ACM

Y. Ning, S. Muthiah, H. Rangwala, and N. Ramakrishnan. Modeling precursors for event forecasting via nested multi-instance learning. In *Proceedings of the 22Nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '16, pages 1095–1104, New York, NY, USA, 2016. ACM

S. Muthiah, B. Huang, J. Arredondo, et al. Planned Protest Modeling in News and Social Media. In *AAAI Conference on Artificial Intelligence, January 25-30, 2015*, pages 3920–3927, 2015

J. Schlachter, A. Ruvinsky, L. A. Reynoso, S. Muthiah, and N. Ramakrishnan. Leveraging topic models to develop metrics for evaluating the quality of narrative threads extracted from news stories. *Procedia Manufacturing*, 3:4028–4035, 2015

Y. Ning, S. Muthiah, R. Tandon, and N. Ramakrishnan. Uncovering news-twitter reciprocity via interaction patterns. In *Advances in Social Networks Analysis and Mining (ASONAM), 2015 IEEE/ACM International Conference on*, pages 1–8. IEEE, 2015

N. Ramakrishnan, P. Butler, S. Muthiah, N. Self, et al. ‘Beating the news’ with EMBERS: Forecasting Civil Unrest using Open Source Indicators. In *International Conference on Knowledge Discovery and Data Mining, KDD, August 24 - 27, 2014*, pages 1799–1808, 2014

A. Doyle, G. Katz, K. Summers, C. Ackermann, I. Zavorin, Z. Lim, S. Muthiah, P. Butler, N. Self, L. Zhao, et al. Forecasting significant societal events using the embers streaming predictive analytics system. *Big Data*, 2(4):185–195, 2014

Y. Ning, S. Muthiah, N. Ramakrishnan, H. Rangwala, and D. Mares. When do crowds turn violent? uncovering triggers from media. In *2018 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining (ASONAM)*. IEEE, 2018

M. R. Islam, S. Muthiah, B. Adhikari, B. A. Prakash, and N. Ramakrishnan. Deepdiffuse: Predicting the ‘who’ and ‘when’ in cascades. In *2018 IEEE International Conference on Data Mining (ICDM)*, Nov 2018

ACTIVITIES

- **Intellectual Property:** P. Butler, S. Muthiah, and N. Ramakrishnan. Alert Generation from Multiple Streams, 2016. VTIP 17-047
- **Program Committee:** ACM SIGKDD 2018-19, WWW 2019, ASONAM 2018

AWARDS AND HONORS

- **Student Travel Award** 2016
KDD
- **Pratt Fellowship** 2015
Virginia Tech
- **Deployed Application Paper Award** 2014
IAAI

TECHNICAL SKILLS

- **Programming:** Python, C/C++, R, Java, Matlab, HTML, PHP
- **Frameworks:** NOSQL: MongoDB, ElasticSearch, Tensorflow, PyTorch MapReduce

PARTICIPATION IN CONFERENCES

- **KDD 2016**, 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, San Francisco, USA, August 2016.
- **KDD 2015**, 21st ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, Sydney, Australia, August 2015.
- **IAAI 2014**, 27th Conference on Innovative Applications of Artificial Intelligence, Austin, USA, Jan 2014.